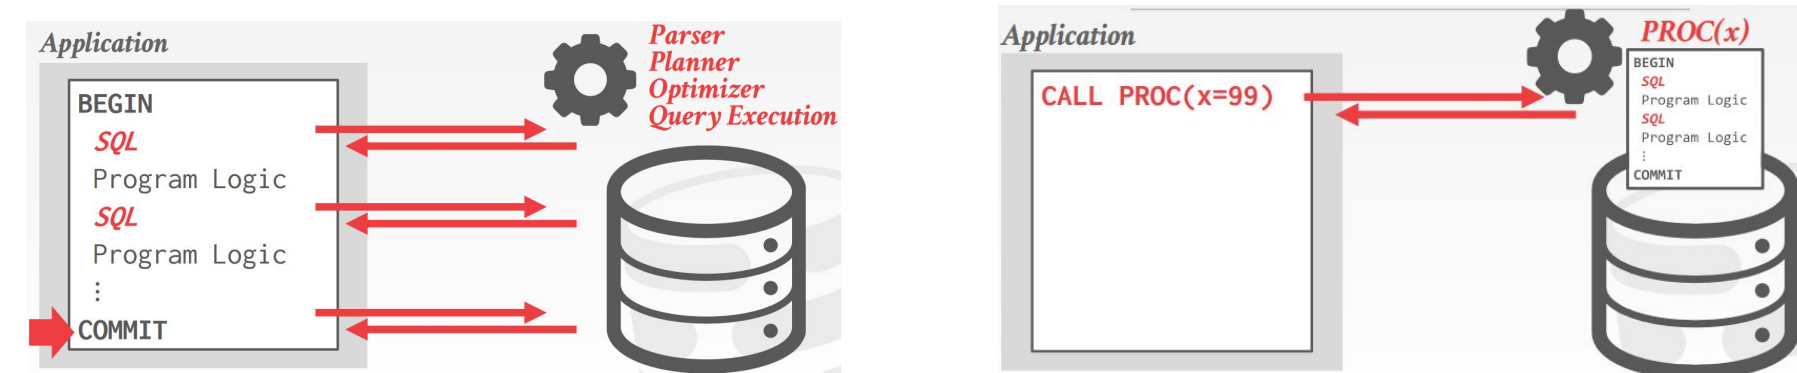


Introduction and Background

- DBMSs provide a way to implement server-side logic by writing UDFs.
 - Mixed imperative and relational.

```
FUNCTION EuroToDollar(money):
  DECLARE rate = SELECT rate FROM rates
  WHERE origin='euro'
  AND dest='dollar';
  RETURN rate * money;
```

- Reduced Overhead.

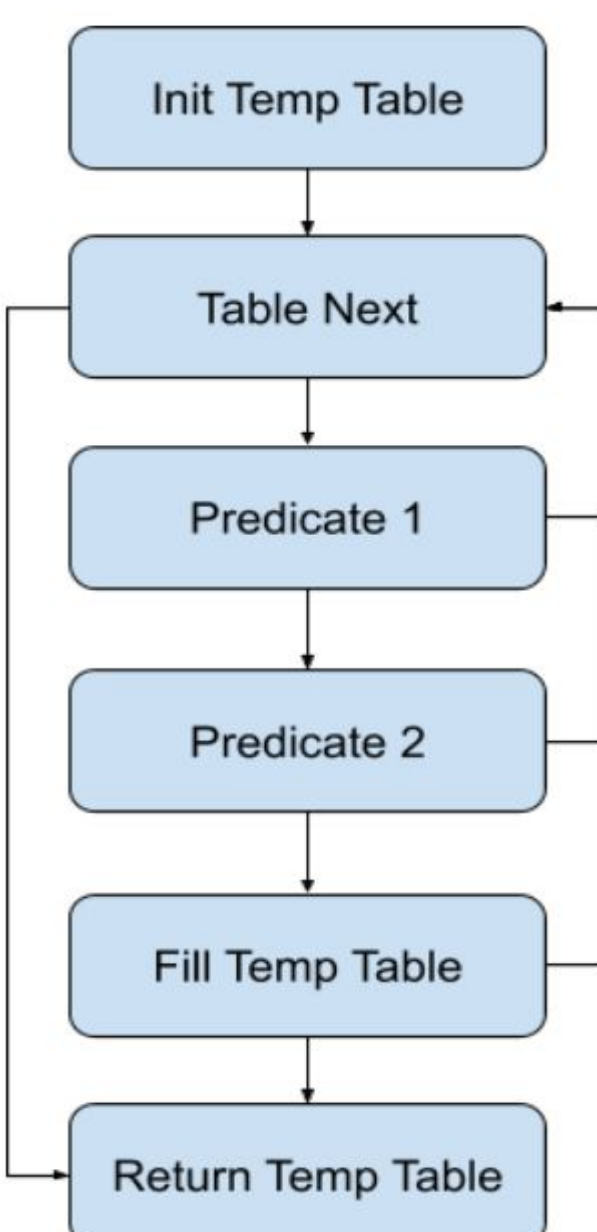


- Traditional UDF execution context switches between imperative and relational execution.
 - Previous work found that context switching causes significant overhead.
- Recent solutions: Imperative to relational conversion, optimize, then execute.
 - Currently limited to scalar-valued functions.
- Our approach: Convert UDF custom high level IR (SQLIR), optimize, then execute.

SQLIR

```
func @Select(%arg0: i64) -> i64 attributes {!isSelect = true} {
  %0 = sqlir.NewTempTable : () -> i64
  br ^bb1
^bb1: // 4 preds: ^bb0, ^bb2, ^bb3, ^bb4
  %c37_i64 = constant 37 : i64
  %1 = sqlir.TableNext %c37_i64 : (i64) -> i1
  cond br %1, ^bb2, ^bb5
^bb2: // pred: ^bb1
  %c37_i64_0 = constant 37 : i64
  %c1_i64 = constant 1 : i64
  %2 = sqlir.GetColumn %c37_i64_0, %c1_i64 : (i64, i64) -> i64
  %c37_i64_1 = constant 37 : i64
  %c2_i64 = constant 2 : i64
  %3 = sqlir.GetColumn %c37_i64_1, %c2_i64 : (i64, i64) -> i64
  %4 = cmpi "slt", %2, %3 : i64
  cond br %4, ^bb3, ^bb1
^bb3: // pred: ^bb2
  %c37_i64_2 = constant 37 : i64
  %c73_i64 = constant 73 : i64
  %5 = cmpi "sgt", %c37_i64_2, %c73_i64 : i64
  cond br %5, ^bb4, ^bb1
^bb4: // pred: ^bb3
  %c37_i64_3 = constant 37 : i64
  %c37_i64_4 = constant 37 : i64
  %6 = muli %c37_i64_3, %c37_i64_4 : i64
  sqlir.FillResult %6, %0 : (i64, i64) -> ()
  %c37_i64_5 = constant 37 : i64
  %c1_i64_6 = constant 1 : i64
  %7 = sqlir.GetColumn %c37_i64_5, %c1_i64_6 : (i64, i64) -> i64
  %c37_i64_7 = constant 37 : i64
  %c2_i64_8 = constant 2 : i64
  %8 = sqlir.GetColumn %c37_i64_7, %c2_i64_8 : (i64, i64) -> i64
  %9 = muli %7, %8 : i64
  sqlir.FillResult %9, %0 : (i64, i64) -> ()
  br ^bb1
^bb5: // pred: ^bb1
  return %0 : i64
}
```

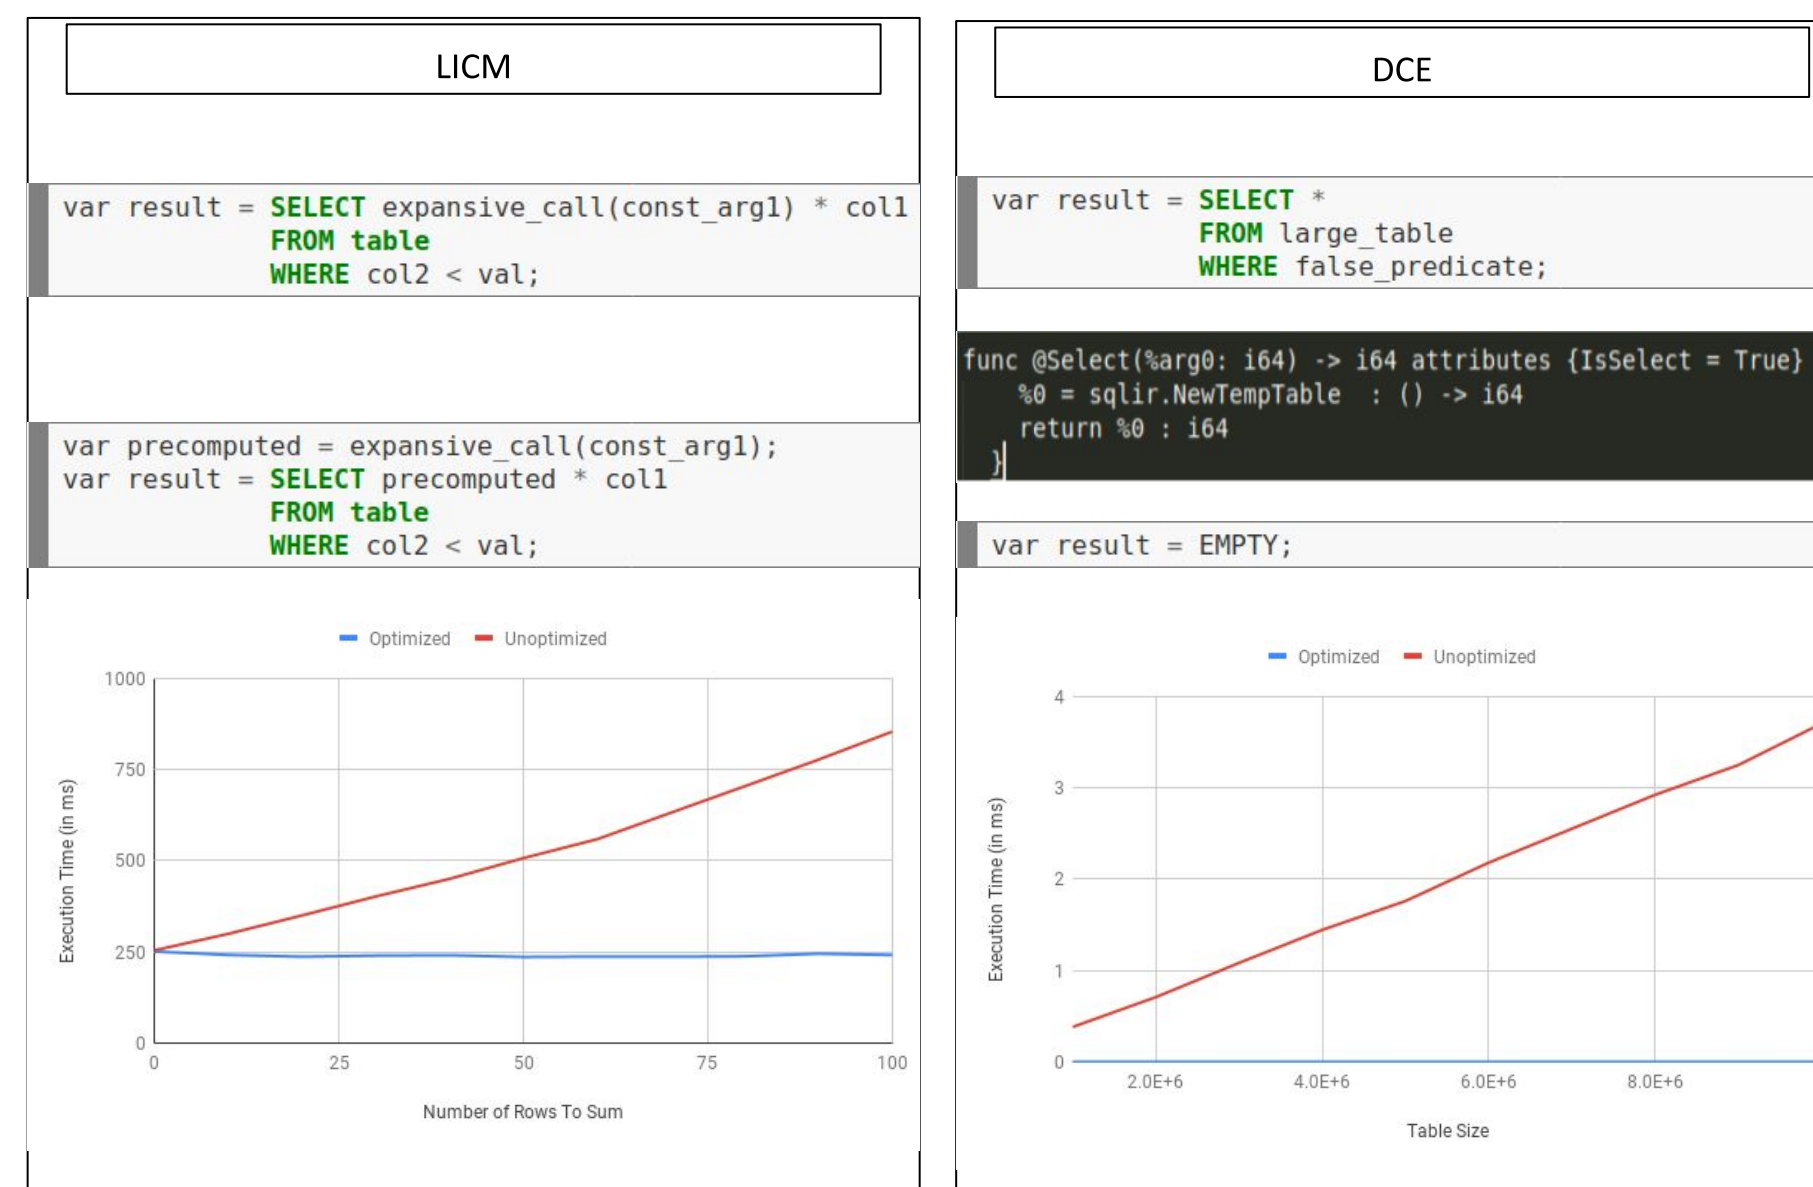
```
SELECT 37*37, col1*col2
FROM F00 WHERE
col1 < col2
AND
37 < 73
```



```
func @Main(%arg0: i64) -> i64 attributes {!isSelect = false} {
  %cst = constant 3.777000e+01 : f64
  %c37_i64 = constant 37 : i64
  %0 = call @Select(%c37_i64) : (i64) -> i64
  %1 = call @Join() : () -> i64
  %c0_i64 = constant 0 : i64
  %c0_i64_0 = constant 0 : i64
  %2 = sqlir.FetchValue %0, %c0_i64, %c0_i64_0 : (i64, i64, i64) -> i64
  %c0_i64_1 = constant 0 : i64
  %c0_i64_2 = constant 0 : i64
  %3 = sqlir.FetchValue %1, %c0_i64_1, %c0_i64_2 : (i64, i64, i64) -> i64
  %4 = muli %2, %3 : i64
  return %4 : i64
}
```

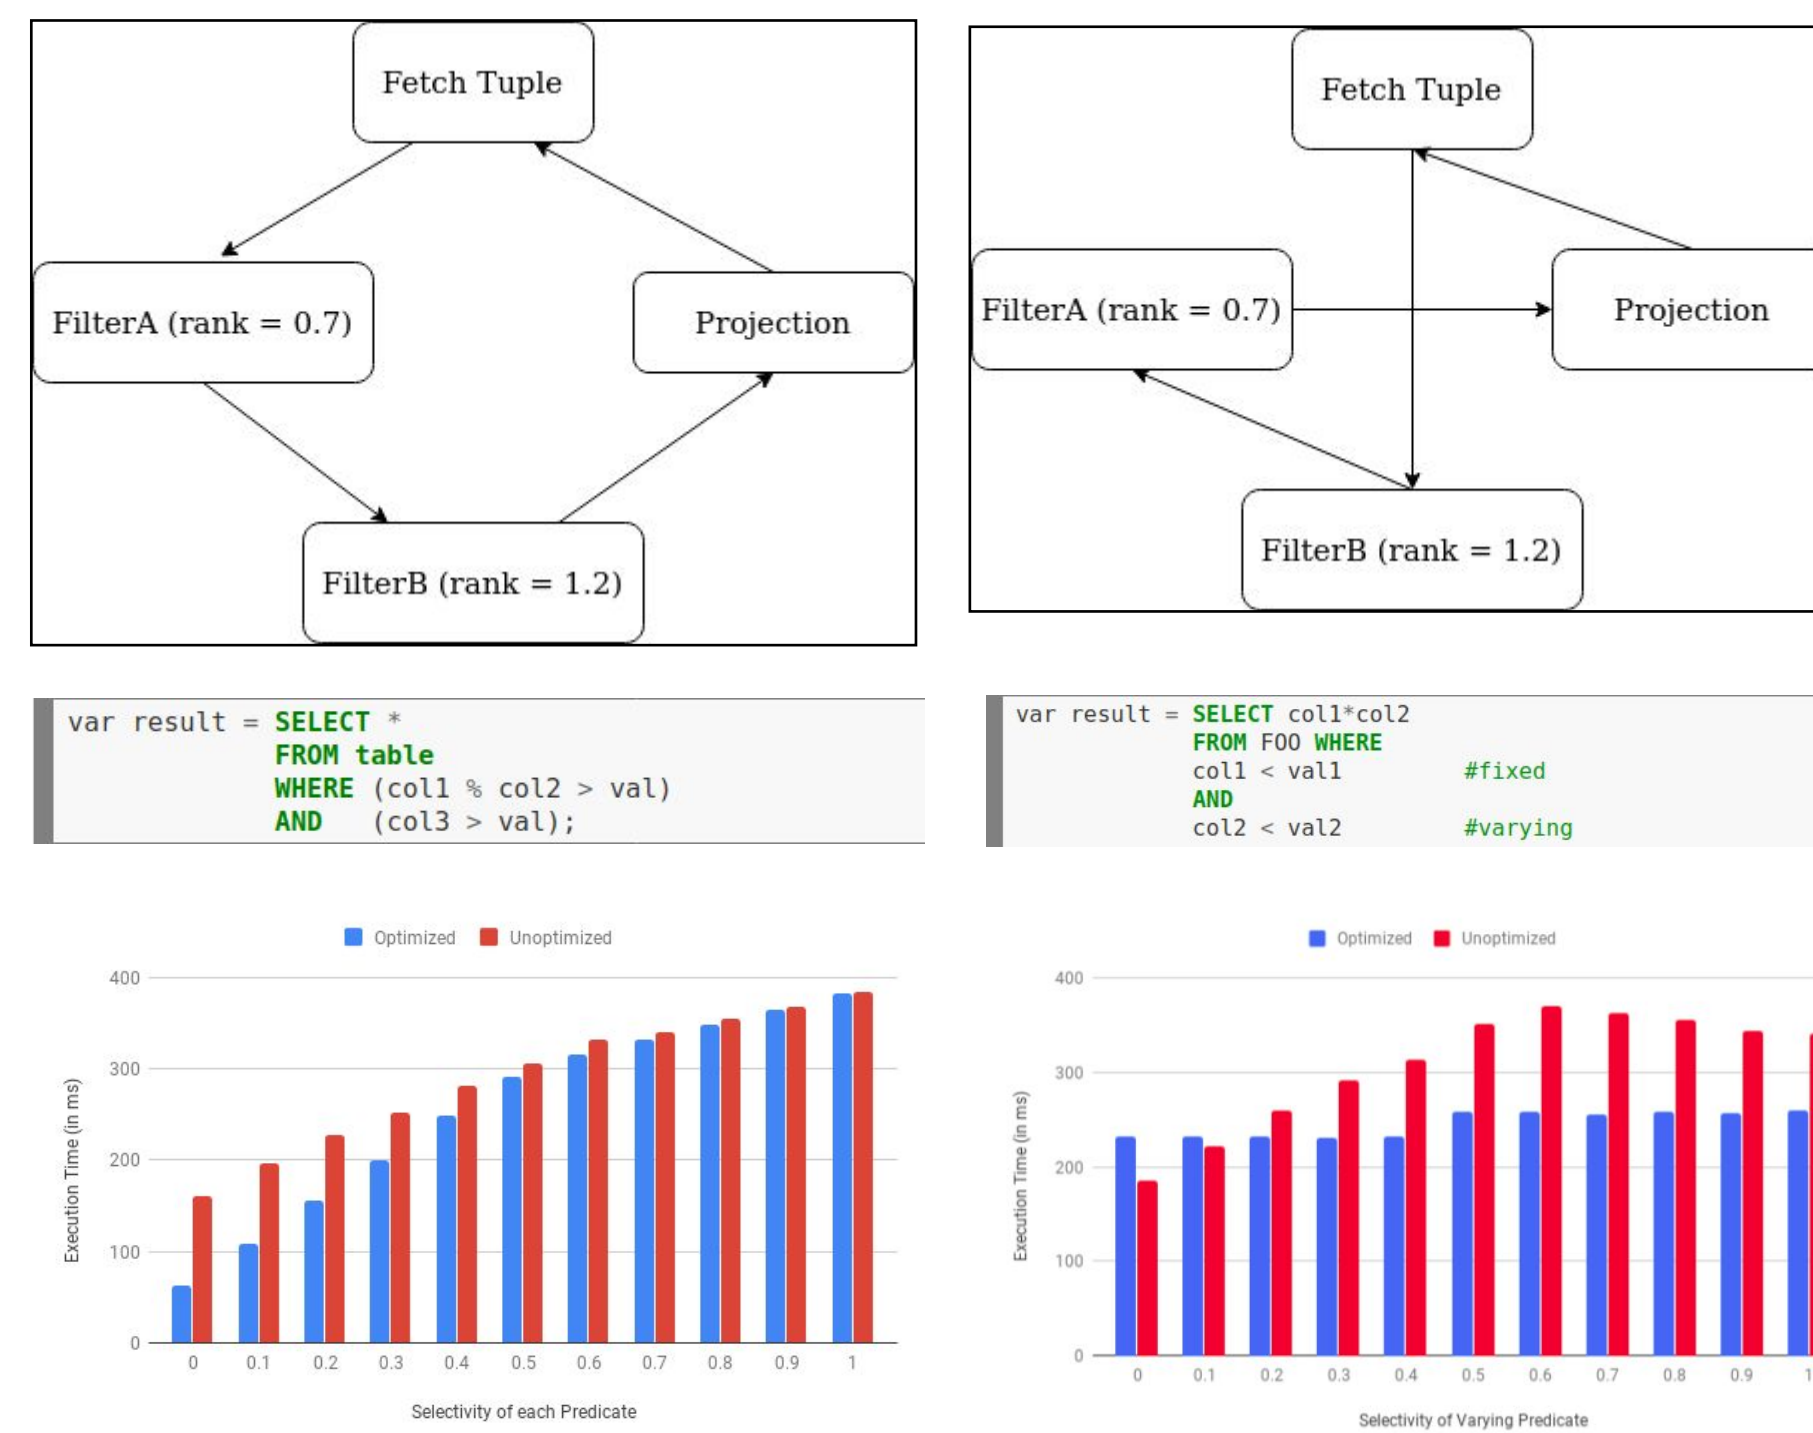
Imperative Optimizations on Selects

- Using MLIR, we get classical imperative optimizations for free.
- Select statements possess statistics on tables.
 - We can perform traditionally unsafe optimization.




Filter Ordering

- Given the structure of our IR, we can easily order filters by changing branch targets.
 - Ordered by rank: (1 - selectivity) / cost-per-tuple.
 - To avoid statistics collection, we only implement 1 / cost-per-tuple.

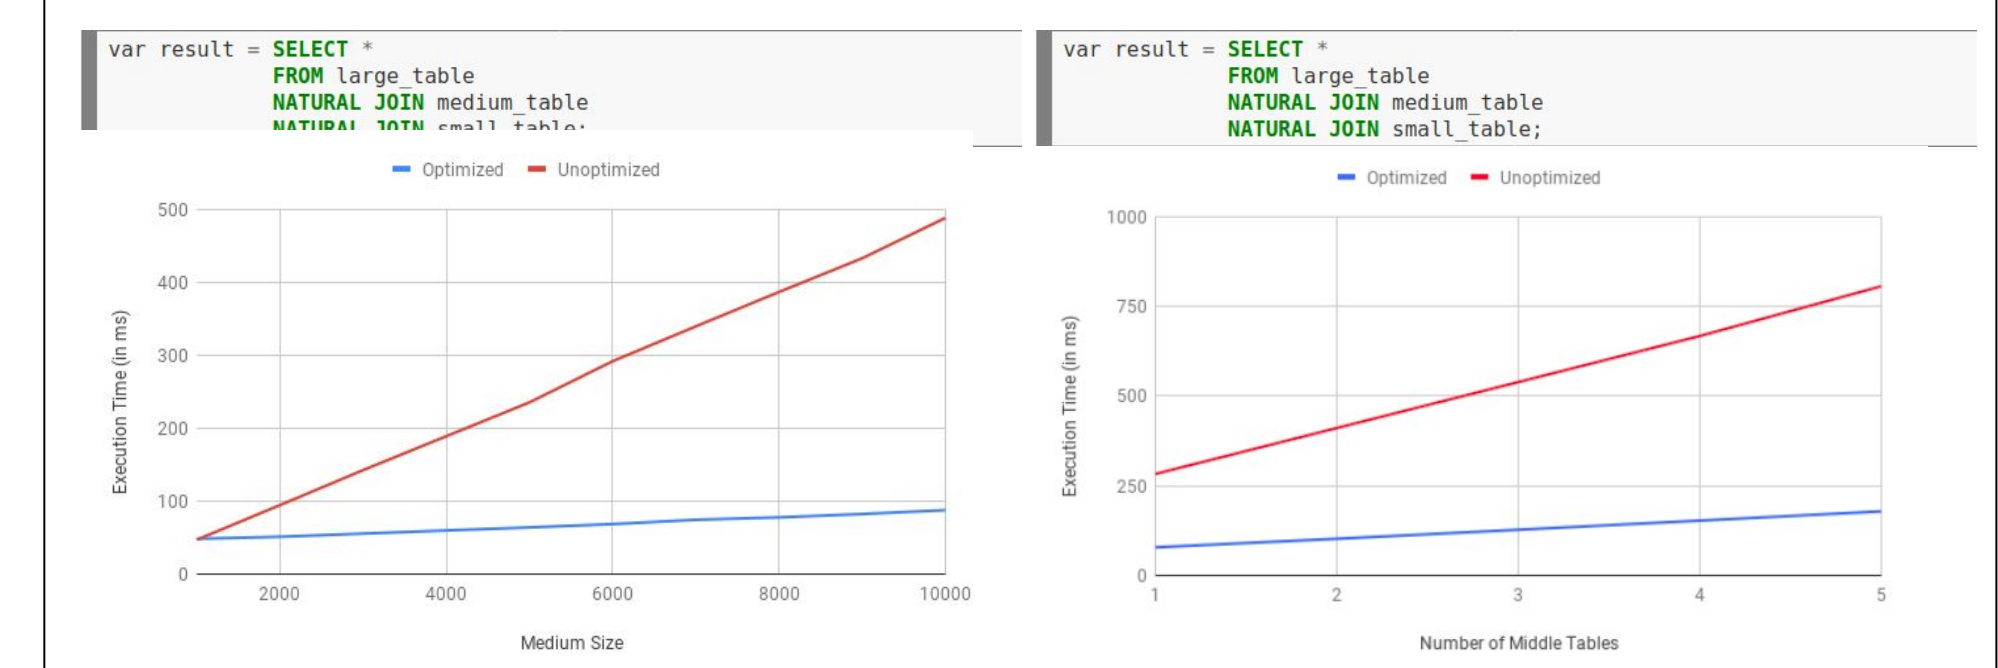



Join Ordering

```
func @Join() -> i64 attributes {!isSelect = false} {
  %c1_i64 = constant 1 : i64
  %c2_i64 = constant 2 : i64
  %0 = sqlir.Join %c1_i64, %c2_i64 : (i64, i64) -> i64
  %c3_i64 = constant 3 : i64
  %1 = sqlir.Join %c3_i64, %0 : (i64, i64) -> i64
  %c4_i64 = constant 4 : i64
  %2 = sqlir.Join %c4_i64, %1 : (i64, i64) -> i64
  %c5_i64 = constant 5 : i64
  %3 = sqlir.Join %c5_i64, %2 : (i64, i64) -> i64
  %c6_i64 = constant 6 : i64
  %4 = sqlir.Join %c6_i64, %3 : (i64, i64) -> i64
  %c7_i64 = constant 7 : i64
  %5 = sqlir.Join %c7_i64, %4 : (i64, i64) -> i64
  return %5 : i64
}
```



```
func @Join() -> i64 attributes {!isSelect = false} {
  %c1_i64 = constant 1 : i64
  %c2_i64 = constant 2 : i64
  %c3_i64 = constant 3 : i64
  %c4_i64 = constant 4 : i64
  %c5_i64 = constant 5 : i64
  %c6_i64 = constant 6 : i64
  %c7_i64 = constant 7 : i64
  %0 = sqlir.Join %c2_i64, %c3_i64 : (i64, i64) -> i64
  %1 = sqlir.Join %c7_i64, %0 : (i64, i64) -> i64
  %2 = sqlir.Join %c4_i64, %1 : (i64, i64) -> i64
  %3 = sqlir.Join %c5_i64, %2 : (i64, i64) -> i64
  %4 = sqlir.Join %c6_i64, %3 : (i64, i64) -> i64
  %5 = sqlir.Join %c1_i64, %4 : (i64, i64) -> i64
  return %5 : i64
}
```



Future Work and Conclusion

- We have not implemented all relational operators, and have not tried to efficiently compose them.
- In addition, this approach could be used for other suchs as Dataframes.
- Optimization framework performs optimizations in a general way combining the optimizations found both in the imperative and the declarative world.
- UDFs lie at the intersection of these two worlds and it's important that we are not restricted while optimizing them.
- Able to leverage optimizations that only appear when both these worlds combine like DCE for Select Statements.
- Did not lose any optimizations in the Declarative world - Join Ordering, Filter Ordering.
- Lot of potential for future work.
- Happy that built this from ground up...!!!